


CM2604 Machine Learning

Nature of Data and Dimensionality Reduction

Week 02 | Prasan Yapa

Overview

- Data in ML
- Data Handling
- Data Cleaning
- Features
- Preparing Dataset
- Dimensionality Reduction



Data in ML

Data in ML

- Any unprocessed fact, value, text, sound, or picture that is not being interpreted and analyzed.
- Most important part of all Data Analytics, Machine Learning, Artificial Intelligence etc.
- Without data, we can't train any model and all modern research and automation will go in vain.
- Big Enterprises are spending lots of money just to gather as much certain data as possible.

Data in ML

Data Types In Machine Learning

Different Forms of Data

- Numeric Data
 - If a feature represents a characteristic measured in numbers , it is called a numeric feature.
- Categorical Data
 - A categorical feature is an attribute that can take on one of the limited, and usually fixed number of possible values based on some qualitative property.
- Ordinal Data
 - This denotes a nominal variable with categories falling in an ordered list. Examples include clothing sizes such as small, medium, and large.

Properties of Data

- Volume: Scale of Data.
- Variety: Different forms of data.
- Velocity: Rate of data streaming and generation.
- Value: Meaningfulness of data in terms of information.
- Veracity: Certainty and correctness in data.

Some Facts on Data

- As compared to 2005, 300 times i.e., 40 Zettabytes ($1\text{ZB}=10^{21}$ bytes) of data will be generated by 2020.
- By 2011, the healthcare sector has a data of 161 Billion Gigabytes.
- 400 Million tweets are sent by about 200 million active users per day.
- Each month, more than 4 billion hours of video streaming is done by the users.
- 30 Billion different types of content are shared every month by the user.
- It is reported that about 27% of data is inaccurate and so 1 in 3 business idealists or leaders don't trust the information on which they are making decisions.

Data Handling

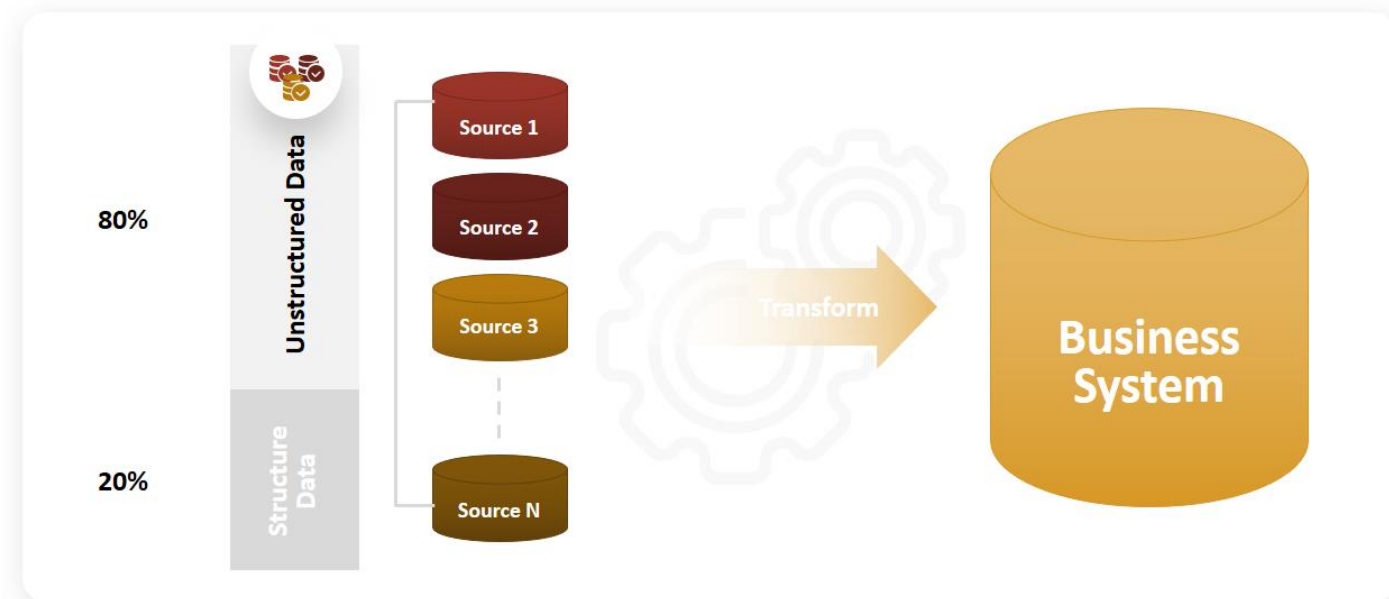
Data Cleaning

- Process of fixing or removing:
 - Incorrect
 - Corrupted
 - Incorrectly formatted
 - Duplicate
 - Incomplete data within a dataset
- If data is incorrect, outcomes and algorithms are unreliable.
- There is no one absolute way to prescribe the exact steps in the data cleaning process.

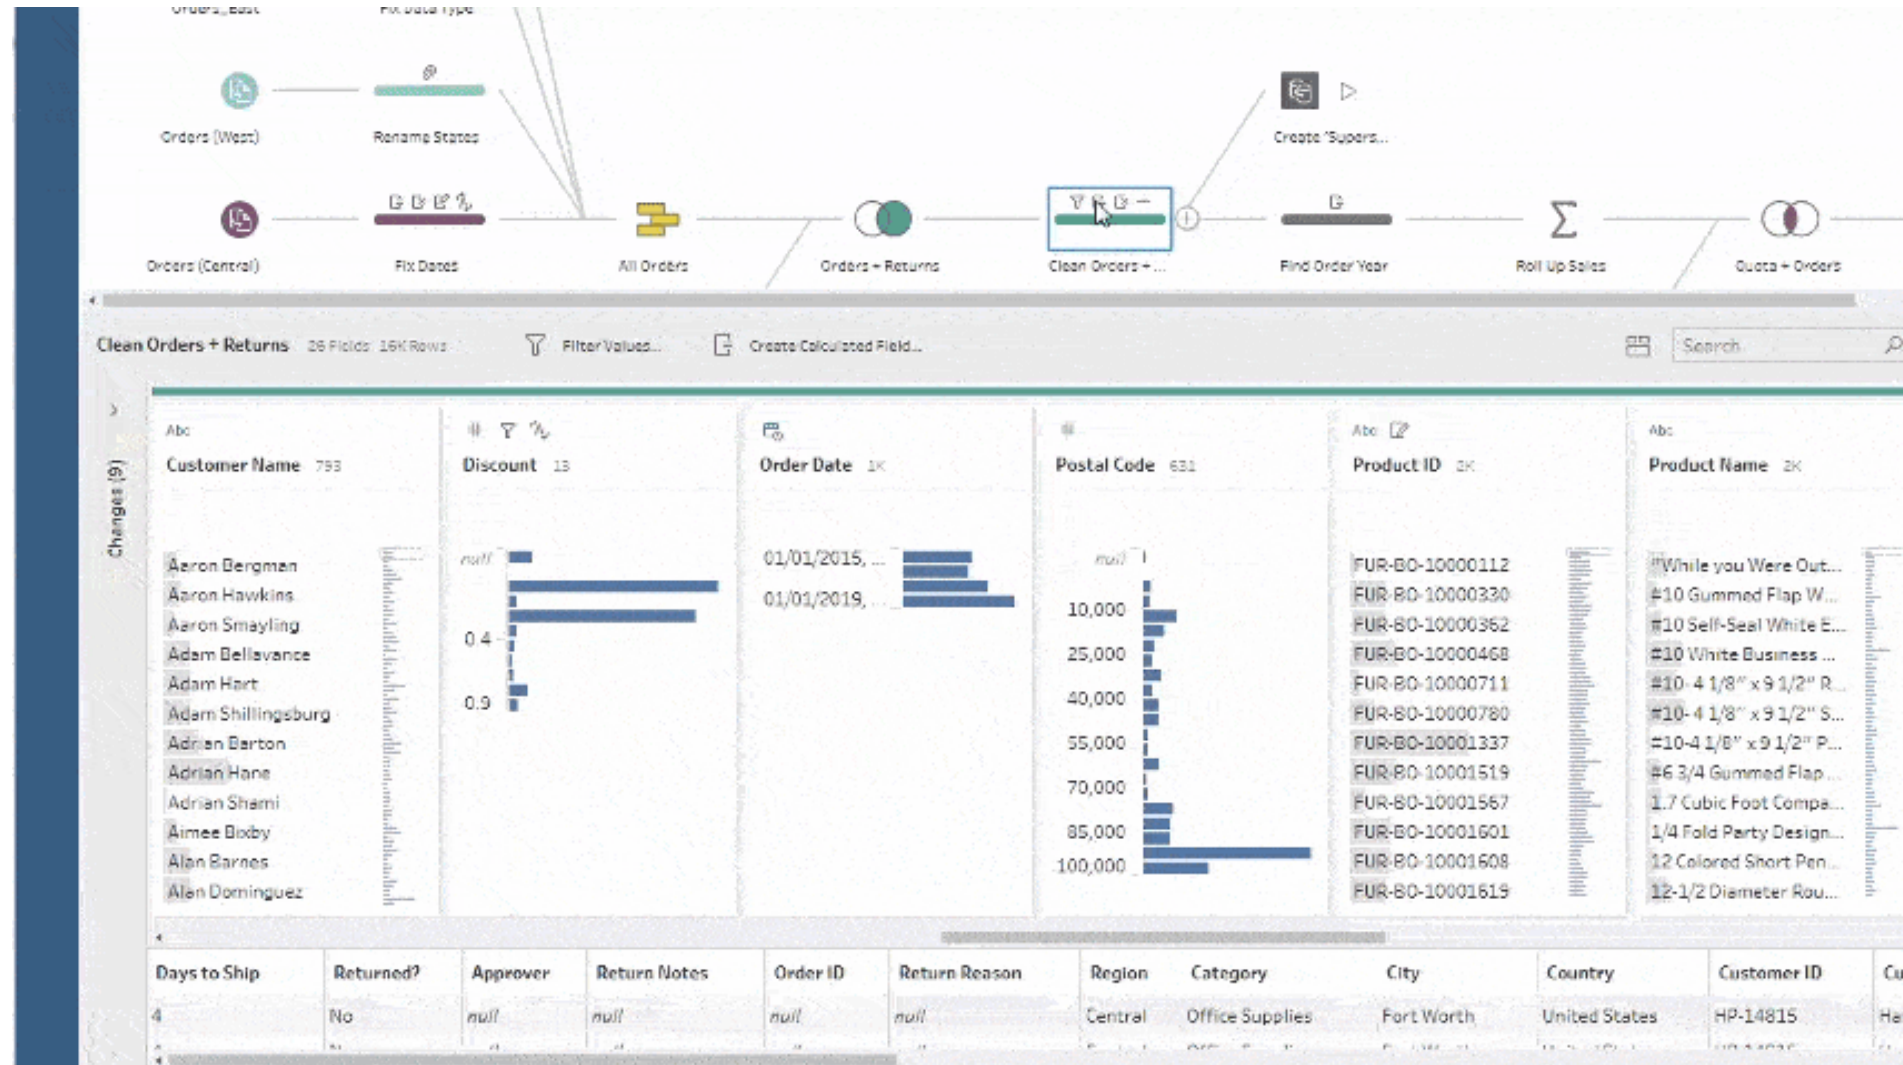
Data Cleaning Vs Data Transformation

- DT is the process of converting data from one format or structure into another.

DATA TRANSFORMATION



How to Clean Data



Step 1: Removing Duplicate or Irrelevant Observations

- Duplicate observations will happen most often during data collection.
- De-duplication is one of the largest areas to be considered in this process.
- Irrelevant observations are when you notice observations that do not fit into the specific problem.
- Creating a more manageable and more performant dataset.

Step 2: Fixing Structural Errors

- Strange naming conventions.
- Typos.
- Incorrect capitalization.
- Inconsistencies can cause mislabeled categories or classes.
 - “N/A” and “Not Applicable” both appear, but they should be analyzed as the same category.

Step 3: Filtering Unwanted Outliers

- There will be one-off observations where do not appear to fit within the data.
- Just because an outlier exists, doesn't mean it is incorrect.
- This task is needed to determine the validity.
- If an outlier proves to be irrelevant for analysis or is a mistake, consider removing it.

Step 4: Handling Missing Data

- You can't ignore missing data because of many algorithms.
- There are a couple of ways to deal with missing data.
 - Observations that have missing values can be dropped, but doing this will drop or lose information, so be mindful of this before removing it.
 - Secondly you can input missing values based on other observations.
 - You might alter the way the data is used to effectively navigate null values.

Step 5: Validation

- Does the data make sense?
- Does the data follow the appropriate rules for its field?
- Does it prove or disprove your working theory, or bring any insight to light?
- Can you find trends in the data to help you form your next theory?
- If not, is that because of a data quality issue?

Stemming Vs Lemmatization in Data Science

- Stemming
 - Process of removing last few characters from a word, often leading to incorrect meanings and spelling.
- Lemmatization
 - Considers the context and converts the word to its meaningful base form, which is called Lemma.

```
1 import nltk
2 from nltk.stem import PorterStemmer
```

```
1 words=['done','doing','studying','identify','this']
2 ps=PorterStemmer()
3 for word in words:
4     print(f"{word}: {ps.stem(word)}")
```

```
done: done
doing: do
studying: studi
identify: identifi
this: thi
```

```
1 from nltk.stem import WordNetLemmatizer
2 lemmatizer = WordNetLemmatizer()
```

```
1 words=['feet','dogs','children','identify','this']
2 for word in words:
3     print(f"{word}: {lemmatizer.lemmatize(word)}")
```


```
feet: foot
dogs: dog
children: child
identify: identify
this: this
```

Variables (Features)




Feature

- A feature is a measurable property of the object you're trying to analyze.
- Features are the basic building blocks of datasets.
- Features are also referred to as “variables” or “attributes.”
- Different business problems do not necessarily require the same features for understanding the business goals.

Activity signals



Feature extraction and Model



Activity prediction



Types of Variables

- Independent Variables Vs Dependent Variables.
- If a variable's value changes when another of the variables change then it is dependent, otherwise independent.
- Independent variables are the input for a process that is being analyzed.
- Dependent variables are the output of the process.

Types of Variables

raw_data - DataFrame

| Index | Age | Education | Income | Marital Status | Purchased |
|-------|-------|-------------|--------|----------------|-----------|
| 0 | 36-55 | Masters | High | Single | 1 |
| 1 | 18-35 | High School | Low | Single | 0 |
| 2 | 36-55 | nan | High | Single | 1 |
| 3 | 18-35 | PhD | Low | nan | 1 |
| 4 | nan | High School | Low | Single | 1 |
| 5 | 55+ | High School | High | Married | 0 |
| 6 | 55+ | High School | nan | Married | 1 |
| 7 | nan | High School | nan | Married | 1 |
| 8 | 55+ | High School | High | Married | 1 |
| 9 | < 18 | Masters | Low | Single | 0 |

Independent variables

Dependent variables

Format Resize ☒ Background color ☒ Column min/max Cancel OK

Preparing Your Dataset

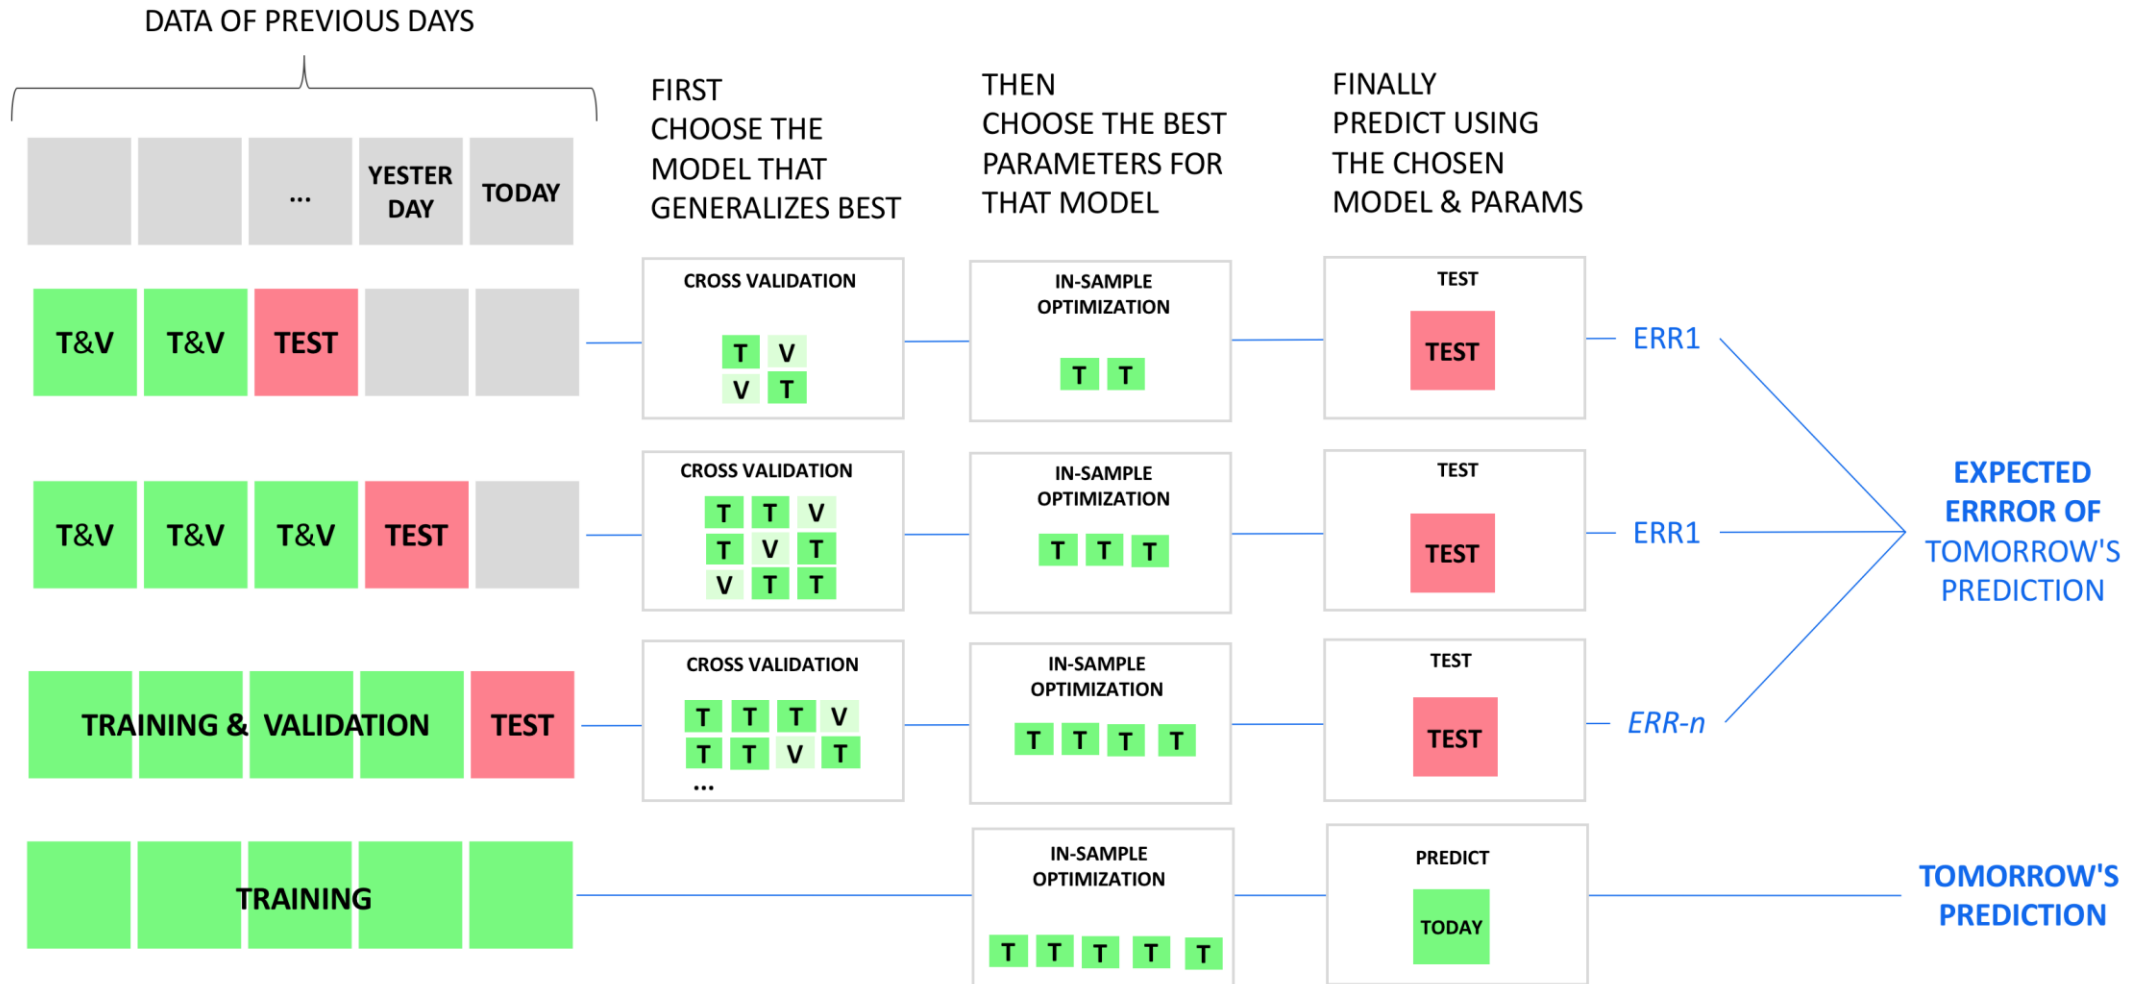
Preparing Your Dataset for ML

- Articulate the problem early.
- Establish data collection mechanisms.
- Check your data quality.
- Format data to make it consistent.
- Reduce data.
- Complete data cleaning.
- Create new features out of existing ones.
- Rescale data.

Training, Validation & Testing Data

- Training data
 - Responsible for building up the machine learning algorithm.
 - The data scientist feeds these data, which corresponds to an output.
- Validation data
 - During training, validation data infuses new data into the model that it hasn't evaluated before.
- Test data
 - After the model is built, testing data once again validates that it can make accurate predictions.
 - If training and validation data include labels, the testing data should be unlabeled.

PREDICT THE NEXT DAY USING PREVIOUS DAYS DATA



Overfitting Vs Underfitting

- **Overfitting**
 - Good performance on the training data, poor generalization to other data.
 - A high accuracy measured on the training set.
- **Underfitting**
 - Poor performance on the training data and poor generalization to other data.
 - Reduces the accuracy and produces unreliable predictions.

Overfitting and Underfitting are two of the main reasons machine learning models have poor performance.

Fixing **overfitting**:

- Simplify the model (fewer parameters)
- Simplify training data (fewer attributes)
- Constrain the model (regularization)
- Use cross-validation
- Use Early stopping
- Build an ensemble
- Gather more data

Fixing **underfitting**:

- More complex model (more parameters)
- Increase number of features
- Feature engineer should help
- Un-constrain the model (no regularization)
- Reduce noise on the data
- Train for longer

Dimensionality Reduction

Dimensionality Reduction

- Reducing the number of random variables by obtaining a set of principal variables.
- Eliminating redundancy and reducing the possibility of the model overfitting.
- Two components:
 - Feature Selection - Smaller subsets of features are chosen by filtering, wrapping or embedding.
 - Feature Extraction - Reducing the number of dimensions in a dataset in order to model variables.

Principal Component Analysis (PCA)

- PCA is used to compress a dataset onto a lower-dimensional feature subspace.
- Feature selection finds a subset of features while PCA produces a smaller new set.
- PCA helps us to identify patterns in data based on the correlation between features.
- PCA aims to find the directions of maximum variance in high-dimensional data.

Applications of PCA in Machine Learning

- PCA is used to visualize multidimensional data.
- It is used to reduce the number of dimensions in healthcare data.
- PCA can help resize an image.
- It can be used in finance to analyze stock data and forecast returns.
- PCA helps to find patterns in the high-dimensional datasets.



How does PCA Work? – Step 1

- Create data by randomly drawing samples.
- Let's start with 2-dimensional data.

```
import numpy as np
import matplotlib.pyplot as plt
import numpy.random as rnd


# Create random 2d data
mu = np.array([10,13])
sigma = np.array([[3.5, -1.8], [-1.8,3.5]])

print("Mu ", mu.shape)
print("Sigma ", sigma.shape)

# Create 1000 samples using mean and sigma
org_data = rnd.multivariate_normal(mu, sigma, size=(1000))
print("Data shape ", org_data.shape)
```

How does PCA Work? – Step 1

- Mean, μ is: $[10, 13]$.
- Covariance matrix, σ is: $\begin{bmatrix} 3.5 & -1.8 \\ -1.8 & 3.5 \end{bmatrix}$.
- Here is a scatter plot of data:




How does PCA Work? – Step 2

- Normalizing data.
- Normalization is done to scale up all features.
- Mean centering is done ensuring the first PC is in the direction of maximum variance.
- Mean centering is done by subtracting mean from all features or channels.

How does PCA Work? – Step 2

```
# Subtract mean from data  
mean = np.mean(org_data, axis= 0)  
print("Mean ", mean.shape)  
mean_data = org_data - mean  
print("Data after subtracting mean ", org_data.shape, "\n")
```



Scatter Plot of Original Data (Left) and Mean Centered Data (Right)

How does PCA Work? – Step 3

- Computing the covariance of all features dimensions.
- Every covariance matrix is symmetric and positive semi-definite.
- It has orthogonal eigen vectors.
- The size of the covariance matrix will be (2 x 2).

```
# Compute covariance matrix
cov = np.cov(mean_data.T)
cov = np.round(cov, 2)
print("Covariance matrix ", cov.shape, "\n")
```

How does PCA Work? – Step 4

- Computing eigen vectors of the covariance matrix.
- The number of eigen vectors will be the same as the number of features.
- Each eigen vector represents a direction of variance.

```
# Perform eigen decomposition of covariance matrix
eig_val, eig_vec = np.linalg.eig(cov)
print("Eigen vectors ", eig_vec)
print("Eigen values ", eig_val, "\n")
```

How does PCA Work? – Step 4

- The eigen vector corresponding to the largest eigen value will give the direction of maximum variance.
- This is the first principal component.
- Then, the eigen vector corresponding to the 2nd largest eigen value will give the direction of the second largest variance.
- This is the second principal component. And, so on.

```
# Sort eigen values and corresponding eigen vectors in descending order
indices = np.arange(0,len(eig_val), 1)
indices = ([x for _,x in sorted(zip(eig_val, indices))])[:,::-1]
eig_val = eig_val[indices]
eig_vec = eig_vec[:,indices]
print("Sorted Eigen vectors ", eig_vec)
print("Sorted Eigen values ", eig_val, "\n")
```

How does PCA Work? – Step 5

- Computing the explained variance and select N components.
- The optimal way is to compute the explained variance of each feature.
- Computing explained variance by dividing the eigen values by the sum of all eigen values.
- Then, we take the cumulative sum of all eigen values.

```
# Get explained variance
sum_eig_val = np.sum(eig_val)
explained_variance = eig_val/ sum_eig_val
print(explained_variance)
cumulative_variance = np.cumsum(explained_variance)
print(cumulative_variance)
```

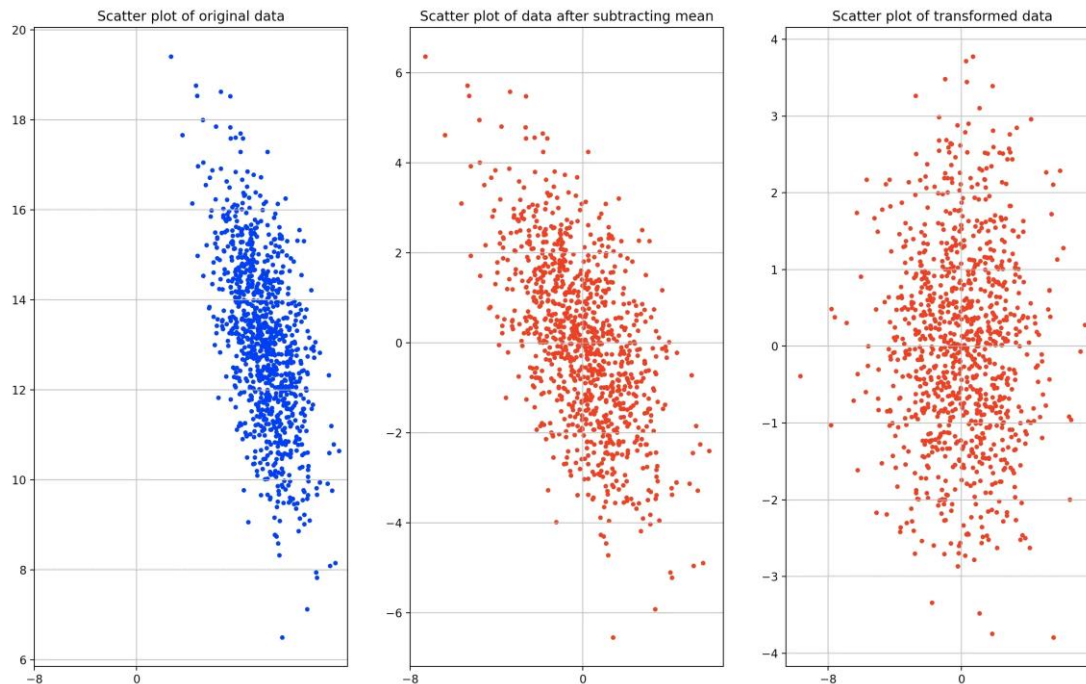
How does PCA Work? – Step 5

- The eigen values here are: [5.50, 1.72].
- The sum of eigen values is: 7.22.
- Explained variance is: [0.76, 0.23].
- Cumulative explained variance is: [0.76, 0.99].
- When we have higher dimensional data, we usually take k components in such a way that we get an explained variance of 0.95 or more.

How does PCA Work? – Step 6

- The dot product of data will be taken with the eigen vectors to get projections of the data in the direction of these eigen vectors.

```
# Take transpose of eigen vectors with data  
pca_data = np.dot(mean_data, eig_vec)  
print("Transformed data ", pca_data.shape)
```

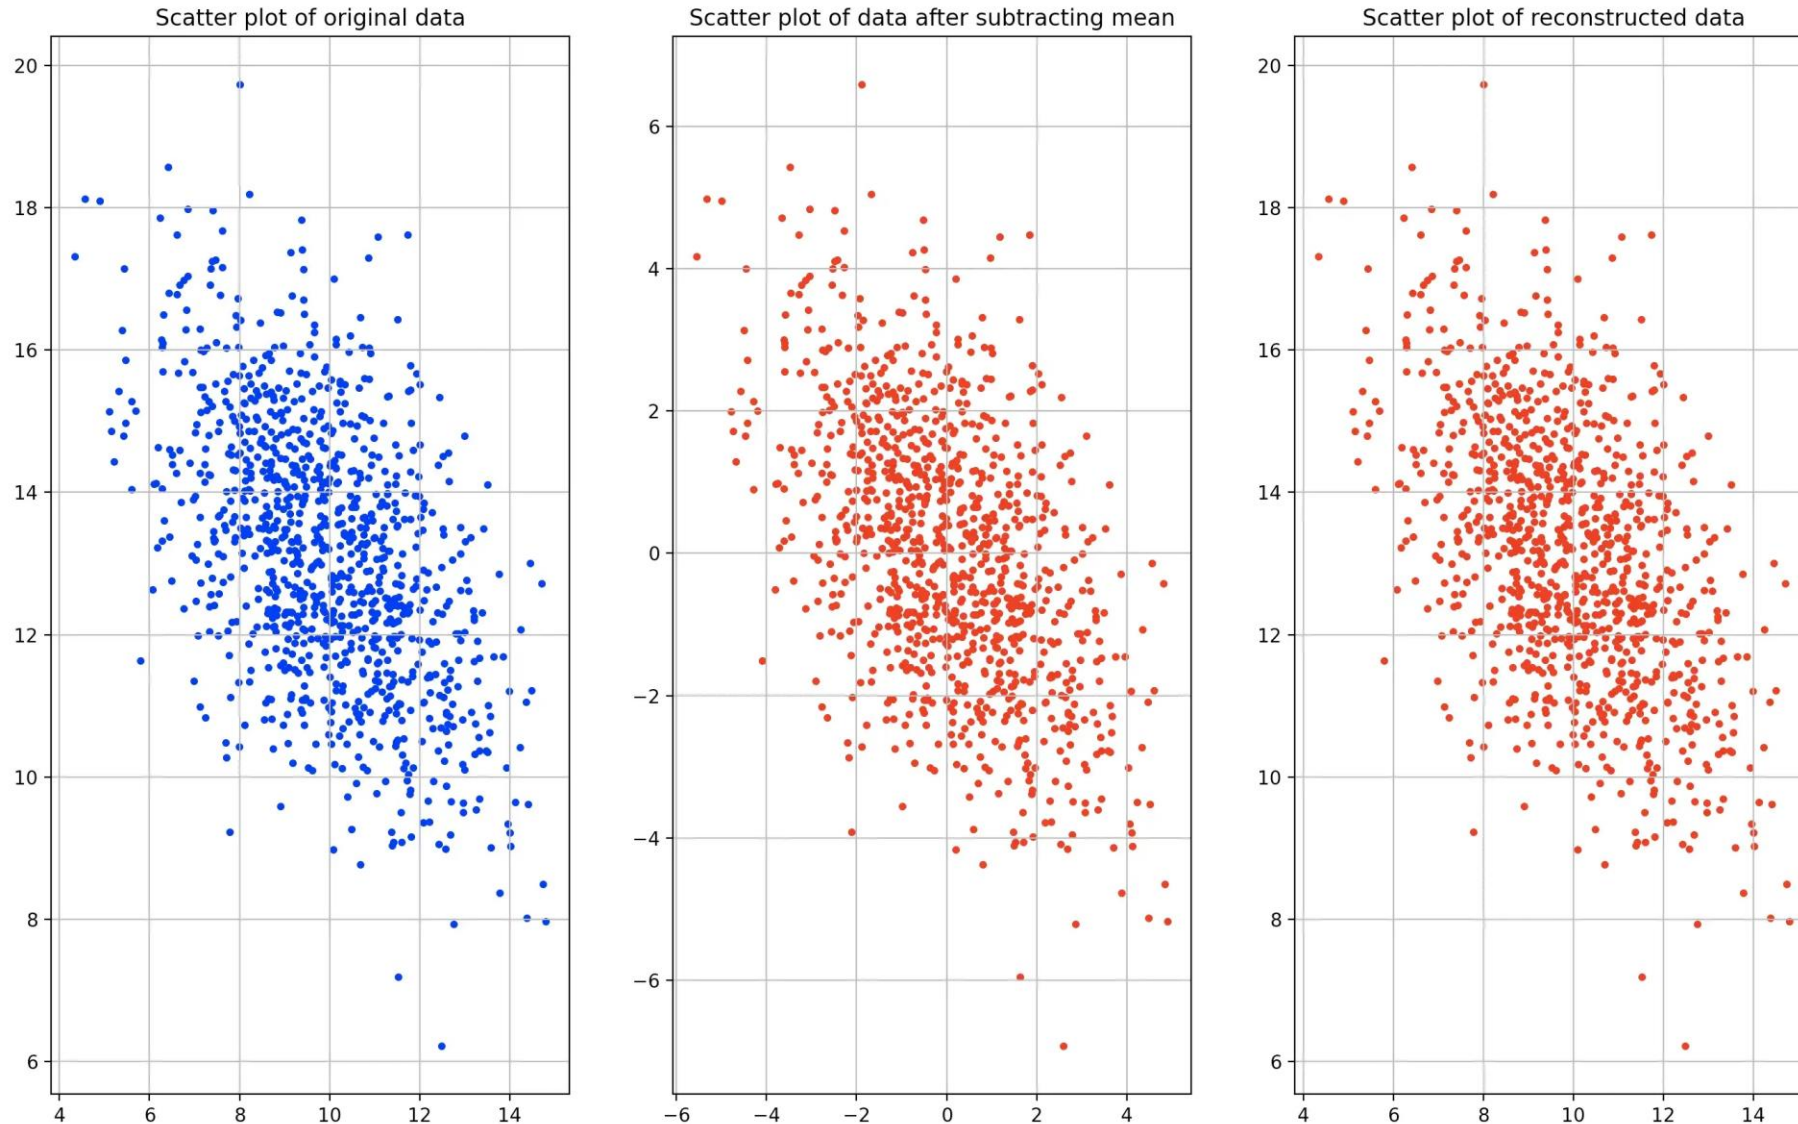


In the scatter plot, we can see that after PCA, the y-axis is the direction of maximum variance.

How does PCA Work? – Step 7

- Inverting PCA and Reconstruct original data.
- Reconstructing the original data by taking the dot product of transpose of eigen vectors.
- Remember to add the mean as well (*mean was subtracted from the data at the beginning to center the data*).
- When we take the dot product of eigen vectors with itself, we get an identity matrix.

How does PCA Work? – Step 7



How does PCA Work? – Step 7

- We can also compute reconstruction loss:

```
# Compute reconstruction loss  
loss = np.mean(np.square(recon_data - org_data))  
print("Reconstruction loss ", loss)
```

- In this case, reconstruction loss is: 2.6426840324903897e-32.
- It is very low because we used all the components to reconstruct the data.

Questions